

Node.js Community Benchmarking Efforts

Gareth Ellis

Node.js Live Paris 2016



About Gareth Ellis

Runtime Performance Analyst @ IBM

Looking at Performance since 2012

Originally solely Java

Started on Node 2015

Member of benchmarking workgroup

Contact me:

gareth.ellis@uk.ibm.com

github.com/gareth-ellis



Agenda

- 1 **Introduction to benchmarking**
 - Key challenges
 - Approaches
 - Identifying source of regression
- **Benchmarking Node.js**
 - Tools
 - Examples
- **Node benchmarking workgroup**
 - Use cases
 - Current benchmarks
 - Results/Graphs
 - How to get Involved



Introduction to benchmarking

- **Change one thing and one thing only between runs**
 - Application code / benchmark
 - Runtime
 - Machine
 - An NPM
- **Performance testing is quite different to functional testing**



Key Challenges

- **Fundamental run-to-run Variance**
 - False positives
 - Collecting enough samples to be sure of the result
 - Documenting Expected Variance
- **Consistent Environment**
 - Known starting machine state
 - Machine Isolation
 - Interleave comparison runs
- **Jumping to Conclusions**



Approaches

- **Micro-Benchmarks**

- Measure a specific function/API
 - Ex: `Buffer.new()`
- Compare key characteristics
- Micro-benchmark improvements may not mean real world improvements
- Risk of not measuring exactly what you expect – especially where a JIT is involved

- **Whole System**

- Benchmark expected customer use case
 - Ex: AcmeAir - <https://github.com/acmeair/acmeair-nodejs>
- The more you test, the more chance for variance

I've found a regression, now what?

- **Are you sure ?**
 - Revalidate environment,
 - Expected variance
- **If so, what changed ?**
 - Your application
 - Node.js
 - Your environment
- **Compare between good/bad cases**
 - Tools
 - Binary search

Benchmarking Node.js

- **Sources of regressions**

- Node.js
 - lib/*.js – buffer, cluster, etc
- V8
- OpenSSL
- libuv
- NPM Module

- **Tools**

- Javascript profiler
 - V8 profiler
 - Appmetrics
- Native profiler (ex perf, tprof, oprofile)



Example – Microbench

```
var harness = require('../../common/harness.js');
```

```
var ARRAY = [1, 2, 23829, 4, 5, 7, 12312321, 2131, 434832, 43792, 23421, 65345, 132210,  
77777, 322131, 1, 2, 23829, 4, 5, 7, 12312321, 2131, 434832, 43792, 23421, 65345, 132210,  
77777, 322131, 1, 2, 23829, 4, 5, 7, 12312321, 2131, 434832, 43792, 23421, 65345, 132210,  
77777, 322131, 1, 2, 23829, 4, 5, 7, 12312321, 2131, 434832, 43792, 23421, 65345, 132210,  
77777, 322131];
```

```
var ITERATIONS = 300000;
```

```
var result;
```

```
function test() {  
    for(var i=0;i<ITERATIONS;i++) {  
        result = new Buffer(ARRAY);  
    }  
}
```

```
harness.run_test(test);
```



Example – Microbench

Node 4.3.2:

`./node benchmark.js`

total time:5.079s / iterations:54 / ops/sec:10.63 / average time:0.09s / variance:0.89%

total time:5.076s / iterations:54 / ops/sec:10.64 / average time:0.09s / variance:0.75%

Node 4.4.0:

`./node benchmark.js`

total time:5.131s / iterations:31 / ops/sec:6.04 / average time:0.17s / variance: 2.32%

total time:5.106s / iterations:31 / ops/sec:6.07 / average time:0.16s / variance: 0.28%

= ~ 40% regression



V8 Profiler

- Part of Node.js binary
- Turn on with
 - `--prof`
- Test-tick-process to post-process
 - `./node --prof-process isolate-0x2818130-v8.log`
- Other helper modules like
 - <https://www.npmjs.com/package/v8-profiler>



Example – v8 profiler (--prof)

```
< 5585 23.7% 23.9% LazyCompile: *fromObject buffer.js:121:20
< 1308 5.6% 5.6% LazyCompile: *subarray native typedarray.js:165:28
< 1263 5.4% 5.4% LazyCompile: *Uint8ArrayConstructByArrayBuffer native typedarray.js:35:42
< 964 4.1% 4.1% Builtin: JSConstructStubGeneric
< 854 3.6% 3.7% Stub: InstanceofStub
< 677 2.9% 2.9% LazyCompile: *test benchmark.js:7:14
< 669 2.8% 2.9% LazyCompile: *Uint8Array native typedarray.js:122:31
---
> 15227 47.1% 47.3% LazyCompile: *fromObject buffer.js:121:20
> 1240 3.8% 3.9% LazyCompile: *subarray native typedarray.js:165:28
> 1166 3.6% 3.6% LazyCompile: *Uint8ArrayConstructByArrayBuffer native typedarray.js:35:42
> 967 3.0% 3.0% Builtin: JSConstructStubGeneric
> 802 2.5% 2.5% Stub: InstanceofStub
> 780 2.4% 2.4% LazyCompile: *test benchmark.js:7:14
> 654 2.0% 2.0% LazyCompile: *Uint8Array native typedarray.js:122:31
```



Perf

```
perf record -i -g -e cycles:u -- ./node --perf-basic-prof benchmark.js
perf report
```

```
diff perf_good.out perf_bad.out
```

```
327c302
```

```
< 91.52% 23.43% node perf-16993.map [...] LazyCompile:*fromObject buffer.js:121
```

```
---
```

```
> 93.25% 46.56% node perf-16934.map [...] LazyCompile:*fromObject buffer.js:121
```

```
331c306
```



--trace-opt --trace-deopt

[marking 0x39570dd44951 <JS Function fromObject (SharedFunctionInfo 0x39570dd12f91)> for recompilation, reason: not much type info but very hot, ICs with typeinfo: 14/64 (21%), generic ICs: 0/64 (0%)]

[compiling method 0x39570dd44951 <JS Function fromObject (SharedFunctionInfo 0x39570dd12f91)> using Crankshaft]

[optimizing 0x39570dd44951 <JS Function fromObject (SharedFunctionInfo 0x39570dd12f91)> - took 0.315, 1.339, 0.511 ms]

[completed optimizing 0x39570dd44951 <JS Function fromObject (SharedFunctionInfo 0x39570dd12f91)>]

[deoptimizing (DEOPT eager): begin 0x39570dd44951 <JS Function fromObject (SharedFunctionInfo 0x39570dd12f91)>

[deoptimizing (eager): end 0x39570dd44951 <JS Function fromObject (SharedFunctionInfo 0x39570dd12f91)>



Binary chop

- Compare changes between good & bad
- Not so bad for adjacent releases
- Bit more difficult for node 0.8 vs node 6....



Appmetrics

- npm install appmetrics
- Can provide cpu, gc, memory, profiling + lots more
- Connect into IBM healthcenter for remote monitoring
- <https://www.npmjs.com/package/appmetrics>

```
var appmetrics = require('appmetrics');
var monitoring = appmetrics.monitor();

monitoring.on('initialized', function (env) {
    env = monitoring.getEnvironment();
    for (var entry in env) {
        console.log(entry + ':' + env[entry]);
    };
});

monitoring.on('cpu', function (cpu) {
    console.log('[' + new Date(cpu.time) + '] CPU: ' + cpu.process);
});
```



7  lib/buffer.js

✚		@@ -185,7 +185,7 @@ function fromString(string, encoding) {
185	185	function fromArrayLike(obj) {
186	186	const length = obj.length;
187	187	const b = allocate(length);
188	188	- for (let i = 0; i < length; i++)
	188	+ for (var i = 0; i < length; i++)
189	189	b[i] = obj[i] & 255;
190	190	return b;
191	191	}
✚		@@ -276,6 +276,7 @@ Buffer.isEncoding = function(encoding) {
276	276	
277	277	
278	278	Buffer.concat = function(list, length) {
	279	+ var i;
279	280	if (!Array.isArray(list))
280	281	throw new TypeError('"list" argument must be an Array of Buffers');
281	282	
✚		@@ -284,15 +285,15 @@ Buffer.concat = function(list, length) {
284	285	
285	286	if (length === undefined) {
286	287	length = 0;
287	287	- for (let i = 0; i < list.length; i++)
	288	+ for (i = 0; i < list.length; i++)
288	289	length += list[i].length;
289	290	} else {
290	291	length = length >>> 0;
291	292	}
292	293	
293	294	var buffer = Buffer.allocUnsafe(length);
294	295	var pos = 0;
295	295	- for (let i = 0; i < list.length; i++) {
	296	+ for (i = 0; i < list.length; i++) {

The result

- Issue in v8 optimiser.
- Will be fixed there once TurboFan becomes default
- github.com/nodejs/node/pull/5819



Node.js Benchmarking Workgroup

- Mandate to track and evangelize performance gains between node releases
- Key goals
 - Define Use Cases
 - Identify Benchmarks
 - Run/Capture results
- 13 current members
- Meetings every month or so

<https://github.com/nodejs/benchmarking>



Benchmarking Use cases

- Back-end API services – Performance over public infrastructure
- Service oriented architectures (SOA) – very low latency
- Microservice-based applications -
- Generating/serving dynamic web page content
 - Express / hapi / koa / react etc
- Single Page Applications with bidirectional communication over WebSockets and/or HTTP/2
- Agents and Data Collectors
- Small scripts
- For these use cases one or more of the following are often important:
 - consistent low latency
 - ability to support high concurrency
 - throughput
 - fast startup/restart/shutdown
 - low resource usage (memory/cpu)

https://github.com/nodejs/benchmarking/blob/master/docs/use_cases.md



Benchmarks – Progress so far...

- Currently Running
 - Startup time
 - Footprint
 - Time to 'require' a module
 - AcmeAir – throughput, response time, footprint measurements
- In progress
 - URL performance
 - Additions from Node.js benchmarks directory



http://benchmarking.nodejs.org



Node.js Benchmarks


Welcome to the Node.js benchmarks page. This page/data is maintained by the [benchmarking working group](#).

The Benchmark working group's purpose is to gain consensus for an agreed set of benchmarks that can be used to:

- track and evangelize performance gains made between Node releases
- avoid performance regressions between releases

Benchmarks are currently run daily and the updated results published on this page in order to provide visibility and to encourage contributors to look for possible regressions after their commits go in.

If you are interested in benchmarks and would like to help out please join the benchmarking workgroup by raising a pull request on the [repo](#) asking to be added, or simply feel free to attend the next meeting or raise/comment on issues.




Charts

acmeair Ops/s(master)

acmeair Ops/s(4.x)

acmeair Ops/s(0.12.x)



How to get involved

- www.github.com/nodejs/benchmarking
- Take a look at what we're running – and the areas we're looking to get benchmarks to cover
- Something missing?
- Something you don't think is quite right?
- Open an issue!

