

Node.js Community Benchmarking Efforts

Gareth Ellis

Node.js Live London 2016

About Gareth Ellis

Runtime Performance Analyst @ IBM

Looking at Performance since 2012

Originally solely Java

Started on Node 2015

Member of benchmarking workgroup

Contact me:

gareth.ellis@uk.ibm.com

github.com/gareth-ellis

Agenda

- **Introduction to benchmarking**
 - Key challenges
 - Approaches
 - Identifying source of regression
- **Benchmarking Node.js**
 - Tools
 - Examples
- **Node benchmarking workgroup**
 - Use cases
 - Current benchmarks
 - Results/Graphs
 - How to get Involved

Introduction to benchmarking

- **Change one thing and one thing only between runs**
 - Application code / benchmark
 - Runtime
 - Machine
 - An NPM
- **Performance testing is quite different to functional testing**

Key Challenges

- **Fundamental run-to-run Variance**
 - False positives
 - Collecting enough samples to be sure of the result
 - Documenting Expected Variance
- **Consistent Environment**
 - Known starting machine state
 - Machine Isolation
 - Interleave comparison runs
- **Jumping to Conclusions**

Approaches

- **Micro-Benchmarks**

- Measure a specific function/API
 - Ex: Buffer.new()
- Compare key characteristics
- Micro-benchmark improvements may not mean real world improvements
- Risk of not measuring exactly what you expect – especially where a JIT is involved

- **Whole System**

- Benchmark expected customer use case
 - Ex: AcmeAir - <http://github.com/acmeair/acmeair-nodejs>
- The more you test, the more chance for variance

I've found a regression, now what?

- **Are you sure ?**
 - Revalidate environment,
 - Expected variance
- **If so, what changed ?**
 - Your application
 - Node.js
 - Your environment
- **Compare between good/bad cases**
 - Tools
 - Binary search

Benchmarking Node.js

- **Sources of regressions**
 - Node.js
 - lib/*.js – buffer, cluster, etc
 - V8
 - OpenSSL
 - libuv
 - NPM Module
- **Tools**
 - Javascript profiler
 - V8 profiler
 - Appmetrics
 - Native profiler (ex perf, tprof, oprofile)

Example – Microbench

```
var harness = require('../common/harness.js');
```

```
var ARRAY = [1, 2, 23829, 4, 5, 7, 12312321, 2131, 434832, 43792, 23421, 65345, 132210,  
77777, 322131, 1, 2, 23829, 4, 5, 7, 12312321, 2131, 434832, 43792, 23421, 65345, 132210,  
77777, 322131, 1, 2, 23829, 4, 5, 7, 12312321, 2131, 434832, 43792, 23421, 65345, 132210,  
77777, 322131, 1, 2, 23829, 4, 5, 7, 12312321, 2131, 434832, 43792, 23421, 65345, 132210,  
77777, 322131];
```

```
var ITERATIONS = 300000;
```

```
var result;
```

```
function test() {  
    for(var i=0;i<ITERATIONS;i++) {  
        result = new Buffer(ARRAY);  
    }  
}
```

```
harness.run_test(test);
```

Example – Microbench

Node 4.3.2:

`./node benchmark.js`

total time:5.079s / iterations:54 / ops/sec:10.63 / average time:0.09s / variance:0.89%

total time:5.076s / iterations:54 / ops/sec:10.64 / average time:0.09s / variance:0.75%

Node 4.4.0:

`./node benchmark.js`

total time:5.131s / iterations:31 / ops/sec:6.04 / average time:0.17s / variance: 2.32%

total time:5.106s / iterations:31 / ops/sec:6.07 / average time:0.16s / variance: 0.28%

= ~ 40% regression

V8 Profiler

- Part of Node.js binary
- Turn on with
 - `--prof`
- Test-tick-process to post-process
 - `./node --prof-process isolate-0x2818130-v8.log`
- Other helper modules like
 - <https://www.npmjs.com/package/v8-profiler>

Appmetrics

- npm install appmetrics
- Can provide cpu, gc, memory, profiling + lots more
- Connect into IBM healthcenter for remote monitoring
- <https://www.npmjs.com/package/appmetrics>

```
var appmetrics = require('appmetrics');
var monitoring = appmetrics.monitor();
monitoring.on('initialized', function (env) {
    env = monitoring.getEnvironment();
    for (var entry in env) {
        console.log(entry + ':' + env[entry]);
    };
});
monitoring.on('cpu', function (cpu) {
    console.log('[' + new Date(cpu.time) + '] CPU: ' + cpu.process);
});
```

Binary chop

- Compare changes between good & bad
- `git rev-list ^good bad^`
- Bit more time consuming for where there's a lot of change sets
- `git bisect` is quite useful though

Example – v8 profiler (--prof)

```
< 5585 23.7% 23.9% LazyCompile: *fromObject buffer.js:121:20
< 1308 5.6% 5.6% LazyCompile: *subarray native typedarray.js:165:28
< 1263 5.4% 5.4% LazyCompile: *Uint8ArrayConstructByArrayBuffer native typedarray.js:35:42
< 964 4.1% 4.1% Builtin: JSConstructStubGeneric
< 854 3.6% 3.7% Stub: InstanceofStub
< 677 2.9% 2.9% LazyCompile: *test benchmark.js:7:14
< 669 2.8% 2.9% LazyCompile: *Uint8Array native typedarray.js:122:31
---
> 15227 47.1% 47.3% LazyCompile: *fromObject buffer.js:121:20
> 1240 3.8% 3.9% LazyCompile: *subarray native typedarray.js:165:28
> 1166 3.6% 3.6% LazyCompile: *Uint8ArrayConstructByArrayBuffer native typedarray.js:35:42
> 967 3.0% 3.0% Builtin: JSConstructStubGeneric
> 802 2.5% 2.5% Stub: InstanceofStub
> 780 2.4% 2.4% LazyCompile: *test benchmark.js:7:14
> 654 2.0% 2.0% LazyCompile: *Uint8Array native typedarray.js:122:31
```

Example - Perf (system based profiler)

- `perf record -i -g -e cycles:u -- ./node --perf-basic-prof benchmark.js`
- `perf report`

`diff perf_good.out perf_bad.out`

327c302

< 91.52% 23.43% node perf-16993.map [...] LazyCompile:*fromObject buffer.js:121

> 93.25% 46.56% node perf-16934.map [...] LazyCompile:*fromObject buffer.js:121

331c306

--trace-opt --trace-deopt

[marking 0x39570dd44951 <JS Function fromObject (SharedFunctionInfo 0x39570dd12f91)> for recompilation, reason: not much type info but very hot, ICs with typeinfo: 14/64 (21%), generic ICs: 0/64 (0%)]

[compiling method 0x39570dd44951 <JS Function fromObject (SharedFunctionInfo 0x39570dd12f91)> using Crankshaft]


[optimizing 0x39570dd44951 <JS Function fromObject (SharedFunctionInfo 0x39570dd12f91)> - took 0.315, 1.339, 0.511 ms]

[completed optimizing 0x39570dd44951 <JS Function fromObject (SharedFunctionInfo 0x39570dd12f91)>]

[deoptimizing (DEOPT eager): begin 0x39570dd44951 <JS Function fromObject (SharedFunctionInfo 0x39570dd12f91)>]

[deoptimizing (eager): end 0x39570dd44951 <JS Function fromObject (SharedFunctionInfo 0x39570dd12f91)>]

The result

7  lib/buffer.js		
@@ -185,7 +185,7 @@ function fromString(string, encoding) {		
185	185	function fromArrayLike(obj) {
186	186	const length = obj.length;
187	187	const b = allocate(length);
188		- for (let i = 0; i < length; i++)
	188	+ for (var i = 0; i < length; i++)
189	189	b[i] = obj[i] & 255;
190	190	return b;
191	191	}
@@ -276,6 +276,7 @@ Buffer.isEncoding = function(encoding) {		
276	276	
277	277	
278	278	Buffer.concat = function(list, length) {
	279	+ var i;
279	280	if (!Array.isArray(list))
280	281	throw new TypeError('"list" argument must be an Array of Buffers');
281	282	
@@ -284,15 +285,15 @@ Buffer.concat = function(list, length) {		
284	285	
285	286	if (length === undefined) {
286	287	length = 0;
287		- for (let i = 0; i < list.length; i++)
	288	+ for (i = 0; i < list.length; i++)
288	289	length += list[i].length;
289	290	} else {
290	291	length = length >>> 0;
291	292	}
292	293	
293	294	var buffer = Buffer.allocUnsafe(length);
294	295	var pos = 0;
295		- for (let i = 0; i < list.length; i++) {
	296	+ for (i = 0; i < list.length; i++) {

- Issue in v8 optimiser.
- Will be fixed there once TurboFan becomes default
- Until then...
- github.com/nodejs/node/pull/5819

Node.js Benchmarking Workgroup

- Mandate to track and evangelize performance gains between node releases
- Key goals
 - Define Use Cases
 - Identify Benchmarks
 - Run/Capture results
- 13 current members
- Meetings every month or so

<http://github.com/nodejs/benchmarking>
<http://benchmarking.nodejs.org>



Benchmarking Use cases

- Back-end API services
- Service oriented architectures (SOA)
- Microservice-based applications
- Generating/serving dynamic web page content
- Single Page Applications with bidirectional communication over WebSockets and/or HTTP/2
- Agents and Data Collectors
- Small scripts

https://github.com/nodejs/benchmarking/blob/master/docs/use_cases.md

Benchmarking Use cases

- For these use cases one or more of the following are often important:
 - consistent low latency
 - ability to support high concurrency
 - throughput
 - fast startup/restart/shutdown
 - low resource usage (memory/cpu)

https://github.com/nodejs/benchmarking/blob/master/docs/use_cases.md

Benchmarks – Progress so far...

- Currently Running
 - Startup time
 - Footprint
 - Time to 'require' a module
 - AcmeAir – throughput, response time, footprint measurements
 - Docker
- In progress
 - URL performance
 - Additions from Node.js benchmarks directory


Charts

acmeair Ops/s[higher is better](master)

acmeair Ops/s[higher is better](4.x)

acmeair Ops/s[higher is better](0.12.x)

acmeair Ops/s[higher is better](6.x)



How to get involved

- <http://www.github.com/nodejs/benchmarking>
- Take a look at what we're running – and the areas we're looking to get benchmarks to cover
- Something missing?
- Something you don't think is quite right?
- Open an issue!