

Performance Monitoring with Repeatability Tips & Techniques

#DevLin2016

About Gareth

- Worked as a performance analyst at IBM since 2012 until end of September.
- Originally looking at Java (IBM SDK for Java), later also Node.js
- Member of Node.js Community Benchmarking Workgroup
- Moved to Linköping three weeks ago!

Contact me:

twitter.com/gse1986

github.com/gareth-ellis

Agenda

- Introduction to performance testing
- Biggest challenges of performance testing
- Ways to overcome challenges
- Tools to use

Introduction to performance testing

- Quite different to functional & system testing
- Change one thing at a time
 - Application code / benchmark
 - Machine / machine configuration
 - Runtime
- Can't necessarily rely on previous measures
 - Most investigations involve a run of a "good" build and a build under test

Designing a performance test

- What do you want to test?
- Ideal state?
 - Has the system warmed up?
 - Are you going to be measuring something useful?
- What else is needed to ensure a fair test?
 - Saturating server for a throughput test
 - Ensuring that system has adequate resource
- What logs/data should we store?
- How do you tidy up afterwards, so future tests are unaffected?

Methods of benchmarking

- Micro-Benchmarks

- Measure a specific function/API
- Ex: `Buffer.new()`
- Compare key characteristics
- Micro-benchmark improvements may not mean real world improvements
- Risk of not measuring exactly what you expect – especially where a JIT is involved

- Whole System

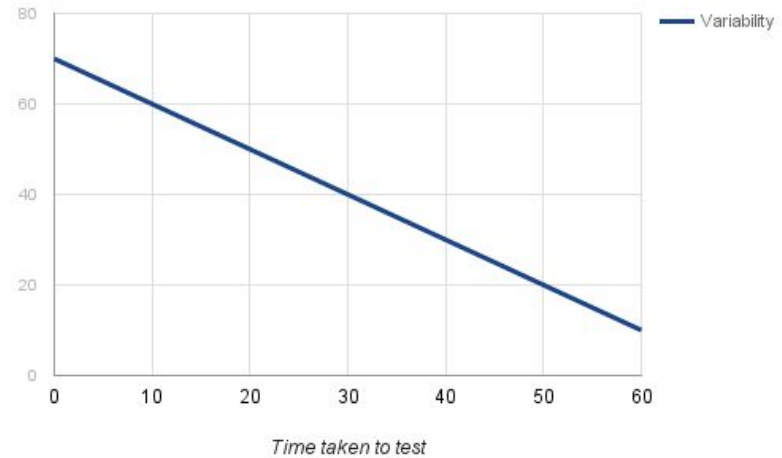
- Benchmark expected customer use case
- Ex: AcmeAir - <http://github.com/acmeair/acmeair-nodejs>
- The more you test, the more chance for variance

Biggest Challenges

- Run-to-run variability
 - False positives
 - Collecting enough samples to be confident in the result
 - Documenting expected variability
 - Repeatability
- Consistent environment
 - Known initial machine state
 - Machine isolation
 - Interleaving runs
- Jumping to conclusions

Reliable data
is typically
expensive....

Variability vs Time



Approaches to overcoming issues

- Setting up the machine for performance testing
 - Turn off power saving
 - Ensure RAM exceeds requirement - paging is bad
 - Ensure the only “bottleneck” is what we’re testing.
- Ensure test approach is repeated to the letter
 - Machine state
 - Software stack state
 - VM state - shared classes

You've found a regression, now what?

- Double check data, is it genuine?
- If so, what changed?
 - Machine?
 - Application code?
 - Software stack?
- Compare between good and bad case
 - Binary search
 - Tooling

Tooling

- Perf - Unix
 - <https://perf.wiki.kernel.org>
- WPT - Windows (as part of Windows SDK)

Thank you - Questions?

<http://github.com/gareth-ellis/slides>